

# GAURAV TADKAPALLY

San Francisco, CA | (213) 913-7899 | [gaurav.tadkapally@usc.edu](mailto:gaurav.tadkapally@usc.edu) | [linkedin.com/in/gauravreddy08](https://www.linkedin.com/in/gauravreddy08) | [gtadkapally.com](https://www.gtadkapally.com)

## EDUCATION

<b>University of Southern California</b>	California, United States
<b>Master of Science in Computer Science: 3.7/4.0</b>	June 2023 - December 2024
- Served as a Teaching Assistant (TA) for the graduate course Applied Machine Learning for Natural Language Processing (ITP 459)	
<b>Vellore Institute of Technology</b>	Andhra Pradesh, India
<b>Bachelor of Technology in Computer Science and Engineering: 8.94/10</b>	May 2019 - May 2023

## EXPERIENCE

<b>Pitney Bowes</b>	Connecticut, United States
<b>Data Science Intern</b>	June 2024-August 2024

- Designed agentic coding assistant for software testing, leveraging **speculative decoding** for accelerated inference speed by 300% and **Abstract Syntax Tree (AST) based retrieval** for document indexing (**Cursor Clone: Demo**)
- Leveraged **Direct Preference Optimization (DPO)** and **LoRA Adapters** for post-training codellama (Llama2), improving model's generative accuracy by 15% (tested via Mutational Testing)
- Implemented retrieval methodologies (SQLite-FTS BM25 & Contextual Embedding) to enhance efficiency and accuracy in retrieving relevant codebase context
- Integrated JaCoCo and Mutational Testing (PIT) to automatically evaluate code coverage & test effectiveness of generated unit tests

<b>MUKHAM</b>	Andhra Pradesh, India
<b>Machine Learning Engineer Intern</b>	October 2022-May 2023

- Optimized facial recognition model for edge deployment (mobile application), leveraging **knowledge distillation, Post-training Quantization (8-bit quantization)**, decreasing model size by 75%
- Designed a Presentation Attack Detection system (facial spoof detection) utilizing the Lucas Kanade algorithm for motion analysis, achieving a 80% success rate in identifying spoofed faces

<b>MUKHAM Pvt Ltd</b>	Andhra Pradesh, India
<b>Research Assistant</b>	October 2022 - May 2023

- Developed a UAV-based wildfire detection algorithm utilizing the EfficientNetB0 architecture, incorporating **Neural Architecture Search (NAS)** for model optimization, resulting in a 98% precision rate
- Engineered smart glasses with an Object Detection model (Incremental Learning) for visually impaired, leading 78% accuracy

## SKILLS AND CERTIFICATIONS

**Languages:** Python, TypeScript, JavaScript  
**ML Stack:** PyTorch, Tensorflow, HuggingFace, LangChain, Agents SDK, TRL, PEFT, Scikit-learn, Pandas, NumPy  
**Tools & Technologies:** AWS (Cloud Practitioner), Azure (AI Fundamentals), SQL, NoSQL, Selenium, Redis

## ACADEMIC PROJECTS

- Poogle: Perplexity Clone (Demo)**
- Engineered a multi-agent web search system with 3 specialized agents, coordinated via shared context memory to decompose tasks, parallelize search, and synthesize high-precision answers
  - Improved token efficiency by 65% via ID-based memory referencing and vector-embedded semantic retrieval, enabling scalable, context-aware web search

- Made AI play Mafia: A multi-agent asynchronous communication (Demo)**
- Developed **asynchronous multi-agent AI system**, enabling structured communication among 6+ autonomous agents in social deduction gameplay scenarios
  - Implemented modular two-part brain architecture (Scheduler & Generator), with a **concurrency-safe** shared context

- AK15: Agentic Kubernetes Middleware (Github)**
- Devised an LLM-based middleware that automates Kubernetes cluster read queries, achieving a 93% reduction in contextual token usage through agentic function calling and context retrieval
  - Implemented 15 specialized API functions enabling the LLM to perform human-like, context-aware interactions with Kubernetes, optimizing and reducing API costs by leveraging targeted data retrieval strategies

## PUBLICATIONS

- Sethuraman, S. C., Reddy Tadkapally, G. et al. **Simplymime: A dynamic gesture recognition and authentication system for smart remote control**. IEEE Sensors Journal (2024). <https://doi.org/10.1109/JSEN.2024.3487070>
- Sethuraman, Sibi C., Gaurav Reddy Tadkapally, et al. **iDrone: IoT-Enabled Unmanned Aerial Vehicles for Detecting Wildfires Using Convolutional Neural Networks**. Springer Nature Computer Science (2022). <https://doi.org/10.1007/s42979-022-01160-7>